

Class 07


A17576411

#Clustering

We will start today's lab with clustering methods, in particular so-called K-means. The main function for this in R is `kmeans()`.

Let's try it on some made up data where we know what the answer should be.

```
x <- rnorm(10000,mean=3)
hist(x)
```




60 points

```
tmp <- c(rnorm(30, mean = 3), rnorm(30, mean = -3))
x <- cbind(x=tmp, y = rev(tmp))
head(x)
```

```
      x      y
[1,] 2.026536 -2.8131156
[2,] 2.292031 -0.9471092
[3,] 4.270506 -2.5665141
[4,] 1.839629 -3.2833284
[5,] 4.133919 -2.9448375
[6,] 2.742538 -4.2206130
```

We can pass this to the base R `plot()` for the function for a quick plot

```
plot(x)
```



```
k <- kmeans(x,centers = 2,nstart = 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

	x	y
1	2.956482	-2.917354
2	-2.917354	2.956482

[illegible]

```
[1] 53.67978 53.67978
      (between_SS / total_SS =  90.6 %)
```

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"   "size"         "iter"         "ifault"
```

k\$size

Q2. Cluster membership?


[illegible]

k\$centers

	x	y
1	2.956482	-2.917354
2	-2.917354	2.956482

3

```
plot(x, col = k$cluster, pch=16)
```



Q5. Cluster the data again into 4 groups

```
k4<-kmeans(x, centers = 4, n = 20)
k4
```

K-means clustering with 4 clusters of sizes 16, 14, 16, 14

Cluster means:

	x	y
1	3.741822	-2.580840
2	2.058951	-3.301941
3	-2.580840	3.741822
4	-3.301941	2.058951

Clustering vector:

```
[1] 2 1 1 2 1 2 2 1 2 1 1 1 2 1 1 1 2 1 2 1 2 2 1 2 2 2 1 1 1 2 4 3 3 3 4 4 4 3
[39] 4 4 3 4 3 4 3 3 3 4 3 3 3 4 3 4 4 3 4 3 3 4
```


Within cluster sum of squares by cluster:

```
[1] 21.58220  7.06899 21.58220  7.06899  
(between_SS / total_SS = 95.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
plot(x, col = k4$cluster, pch=16)
```



K-means is very popular mostly because it is fast and relatively straightforward to run and understand. It has a big limitation in that you need to tell it how many groups (k, or centers) you want.

Hierarchical clustering

The main function in base R is called `hclust()`. You have to pass it in a “distance matrix” not just your input data.

You can generate a distance matrix with the `dist()` function.

```
hc <- hclust( dist(x))
hc
```

Call:

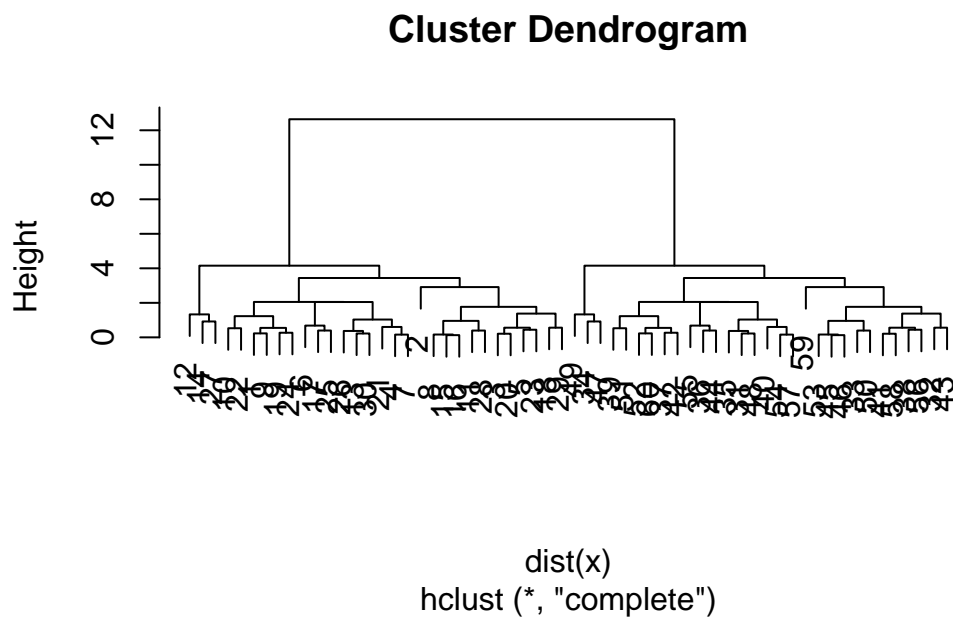
```
hclust(d = dist(x))
```

Cluster method : complete

Distance : euclidean

Number of objects: 60


```
plot(hc)
```



To find the clusters (cluster membership vector) from a `hclust()` result we can “cut” the tree at a certain height that we like.

```
plot(hc)
abline(h=8, col = "red")
```

Cluster Dendrogram




```
grps <- cutree(hc, h=8)
```

```
table(grps)
```

```
grps  
 1  2  
30 30
```

Q6. Plot our hclust results.

```
plot(grps,col=grps, pch=16)
```



#Principal Component Analysis

##PCA of UK food data

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
x
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66	
2	Carcass_meat	245	227	242	267	
3	Other_meat	685	803	750	586	
4	Fish	147	160	122	93	
5	Fats_and_oils	193	235	184	209	
6	Sugars	156	175	147	139	
7	Fresh_potatoes	720	874	566	1033	
8	Fresh_Veg	253	265	171	143	
9	Other_Veg	488	570	418	355	
10	Processed_potatoes	198	203	220	187	
11	Processed_Veg	360	365	337	334	
12	Fresh_fruit	1102	1137	957	674	
13	Cereals	1472	1582	1462	1494	

14	Beverages	57	73	53	47
15	Soft_drinks	1374	1256	1572	1506
16	Alcoholic_drinks	375	475	458	135
17	Confectionery	54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17 5
```

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
dim(x)
```

```
[1] 17 4
```

```
x <- read.csv(url, row.names=1)
head(x)
```

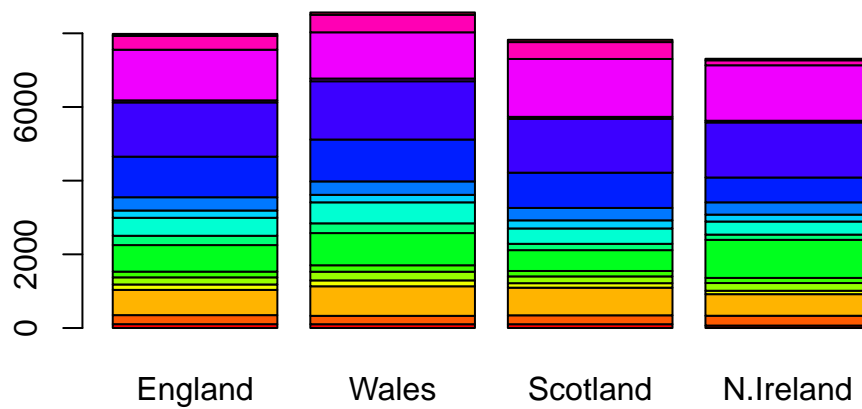
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Personally I prefer the csv rownames -1 one, since if you run the former way more than once, it will continue to remove columns. This for me resulted in me having 3 columns after having accidentally running the program more than once. For this reason also I believe the second way to be more robust.

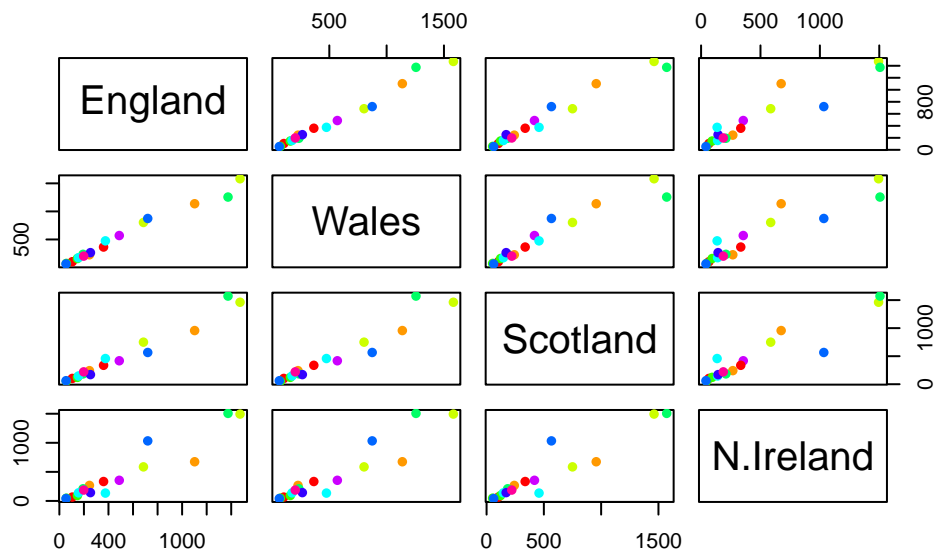
Q3: Changing what optional argument in the above barplot() function results in the following plot?

```
barplot(as.matrix(x), beside=FALSE, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



#I think pairs() graphs two countries against each other. Each of the graphs represents a

Q6. What is the main difference between N. Ireland and the other countries of the UK in terms of this data-set?

A: N. Ireland clearly has different levels on consumption for many foodstuffs, as evidenced by the fact that the line of best fit between N. Ireland and the other countries does not perfectly line up on the diagonal or is not close to it.

Principal Component Analysis (PCA)

PCA can help us make sense of these types of datasets. Let's see how it works.

The main function in base R is called `prcomp()`. In this case we want to first take the transpose of our input `x` so the columns are the food types and the countries are the rows.

```
head(t(x))
```

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175

Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139
	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes		
England	720	253	488		198	
Wales	874	265	570		203	
Scotland	566	171	418		220	
N.Ireland	1033	143	355		187	
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks	
England	360	1102	1472	57	1374	
Wales	365	1137	1582	73	1256	
Scotland	337	957	1462	53	1572	
N.Ireland	334	674	1494	47	1506	
	Alcoholic_drinks	Confectionery				
England	375	54				
Wales	475	64				
Scotland	458	62				
N.Ireland	135	41				

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

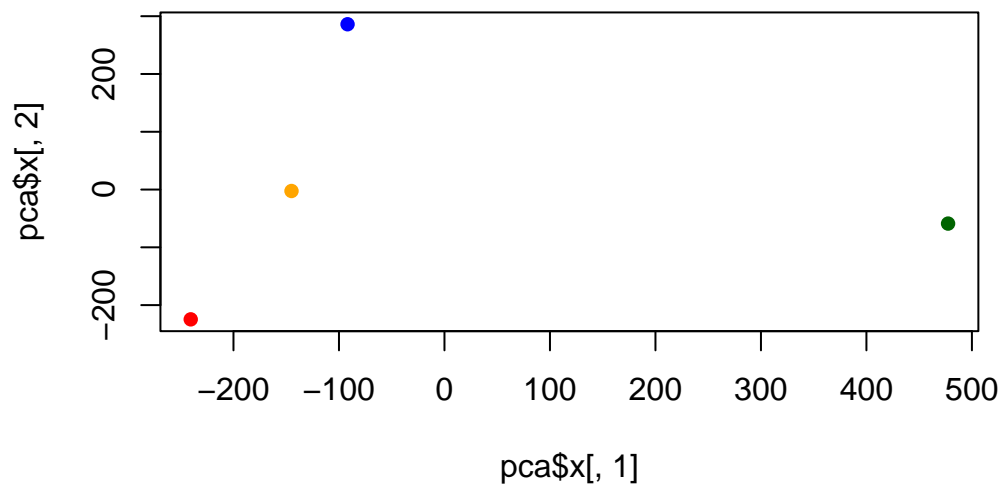
The “loadings” tell us how much the original variables (in our case the foods) contribute to the new variables i.e. the PCs > Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

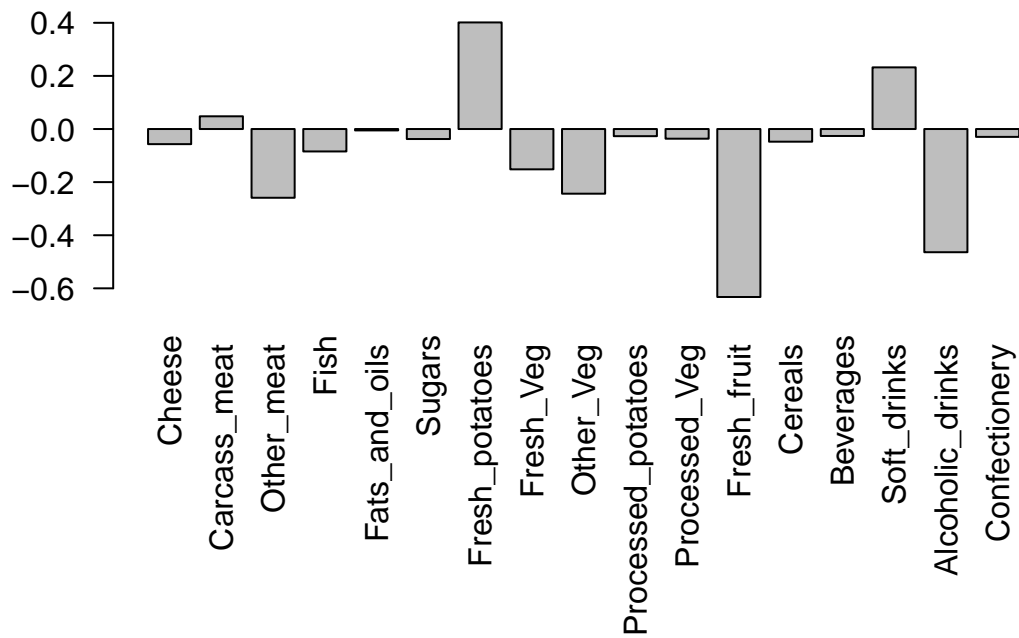
```
plot(pca$x[,1],pca$x[,2], col =c("orange","red","blue","darkgreen"), pch = 16)
```



```
head(pca$rotation)
```


	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.01601285	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.01391582	0.06367111	0.489884628
Other_meat	-0.258916658	-0.01533114	-0.55384854	0.279023718
Fish	-0.084414983	-0.05075495	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.09538866	-0.12522257	0.076097502
Sugars	-0.037620983	-0.04302170	-0.03605745	0.034101334

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Q9. Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```

Soft drinks and fresh potatoes feature prominently. PC2 shows us the rest of the variance not covered in PC1.